

Name: ZHANG XINJIE
Student ID Number: 20M14457
E-mail: zhang.x.az@m.titech.ac.jp

演習 8 勾配降下法

課題

行列 A およびベクトル b が既知のとき、以下の最小二乗法

$$x^* = \arg \min_x \|Ax - b\|_2^2$$

を勾配降下法を用いて未知の x^* を求めるアルゴリズムを実装せよ。初期値は $x = (3, 3)^T$ とし、勾配降下のステップサイズは $\tau = 0.1$ とする。

追加課題

余裕のあるものは、反復ごとの x_1, x_2 を保存しておいて、反復によるこれらの変数の変化をグラフにせよ。さらに、 τ を変えてこのグラフがどうなるか観察せよ。

For nonlinear programming problem, it is difficult to find a optimal solution by finite computation. The optimal solution is expected to be obtained by generating a sequence of points $\{x^{(k)}\}$ that converge to the optimal one.

For a convex function f , the global optimum can be obtained by Gradient Descent Method, which the gradient vector $\nabla f(x^{(k)})$ leads to the direction that the value of the object function will increase or decrease to the optimum point.

Here, the parameter α is proposed as the step to control the speed of converging. By choosing an appropriate starting point $x(0)$, and generating the next point by

$$x^{(k+1)} = x^{(k)} - \alpha^{(k)} \nabla f(x^{(k)})$$

iteratively, the optimal solution can be obtained.

In machine learning, this problem can be considered as linear regression with one variable. Thus, the matrix A must contains one more column of 1 for convenience of computation.

```
In [1]: import numpy as np
import matplotlib.pyplot as plt

a = np.c_[np.asarray([21.3, 22, 26.9, 32.3, 33.1, 38.2])/100]
b = np.c_[np.asarray([116.5, 125.5, 128.1, 132, 141, 145.2])/100]
A = np.concatenate([np.ones([6,1]), a], axis=1)

ATA = np.matmul(A.T, A)
ATb = np.matmul(A.T, b)
# ソルバで解いた解x(デバッグ用)
tx = np.linalg.solve(ATA, ATb)
print(tx)

tau = 0.1 #勾配降下のステップサイズ
#tau == alpha
Niter = 5000 # 反復回数

# xの初期値
x = np.c_[np.asarray([3, 3])]

[[0.88743422]
 [1.47203377]]
```

In this case, the objective function is $f(x) = \|Ax - b\|_2^2$.

```
In [2]: # In machine learning, the function usually is called as cost function or loss function.
def computeCost(A, b, x):
    return sum((A.dot(x)-b)**2)

In [3]: # The parameters are the observed variables, the predictions, the starting point
# step, and number of iterations.
# The parameters that could be used to predict are expected to be obtained.

def gradientDescent(A, b, x, alpha, Niter):
    xtmp = x
    J_history = np.zeros((Niter, 1))
    x_history = np.zeros((Niter, 2, 1))

    for iter in range(Niter):
        #tmp_x0 = x[0] - alpha * sum(A.dot(x)-b)
        tmp_x0 = xtmp[0] - alpha * sum(A.dot(xtmp)-b)
        #tmp_x1 = x[1] - alpha * sum(np.multiply(A.dot(x) - b, A[:,1].reshape(-1,1)))
        tmp_x1 = xtmp[1] - alpha * sum(np.multiply(A.dot(xtmp) - b, A[:,1].reshape(-1,1)))

        xtmp = np.array([tmp_x0, tmp_x1])
        x_history[iter] = xtmp
        #print(xtmp)
        J_history[iter] = computeCost(A, b, xtmp)

    return xtmp, J_history, x_history
```


```
In [4]: x_final, JHistory, xHistory = gradientDescent(A, b, x, tau, Niter)
print(x_final)

[[0.88742668]
 [1.47205971]]
```

The final result of iteration is very close to direct solution, meaning it is correct.

```
In [5]: # Plot the x[0], x[1] changes as the iterations increase.
# They finally converge to a number.

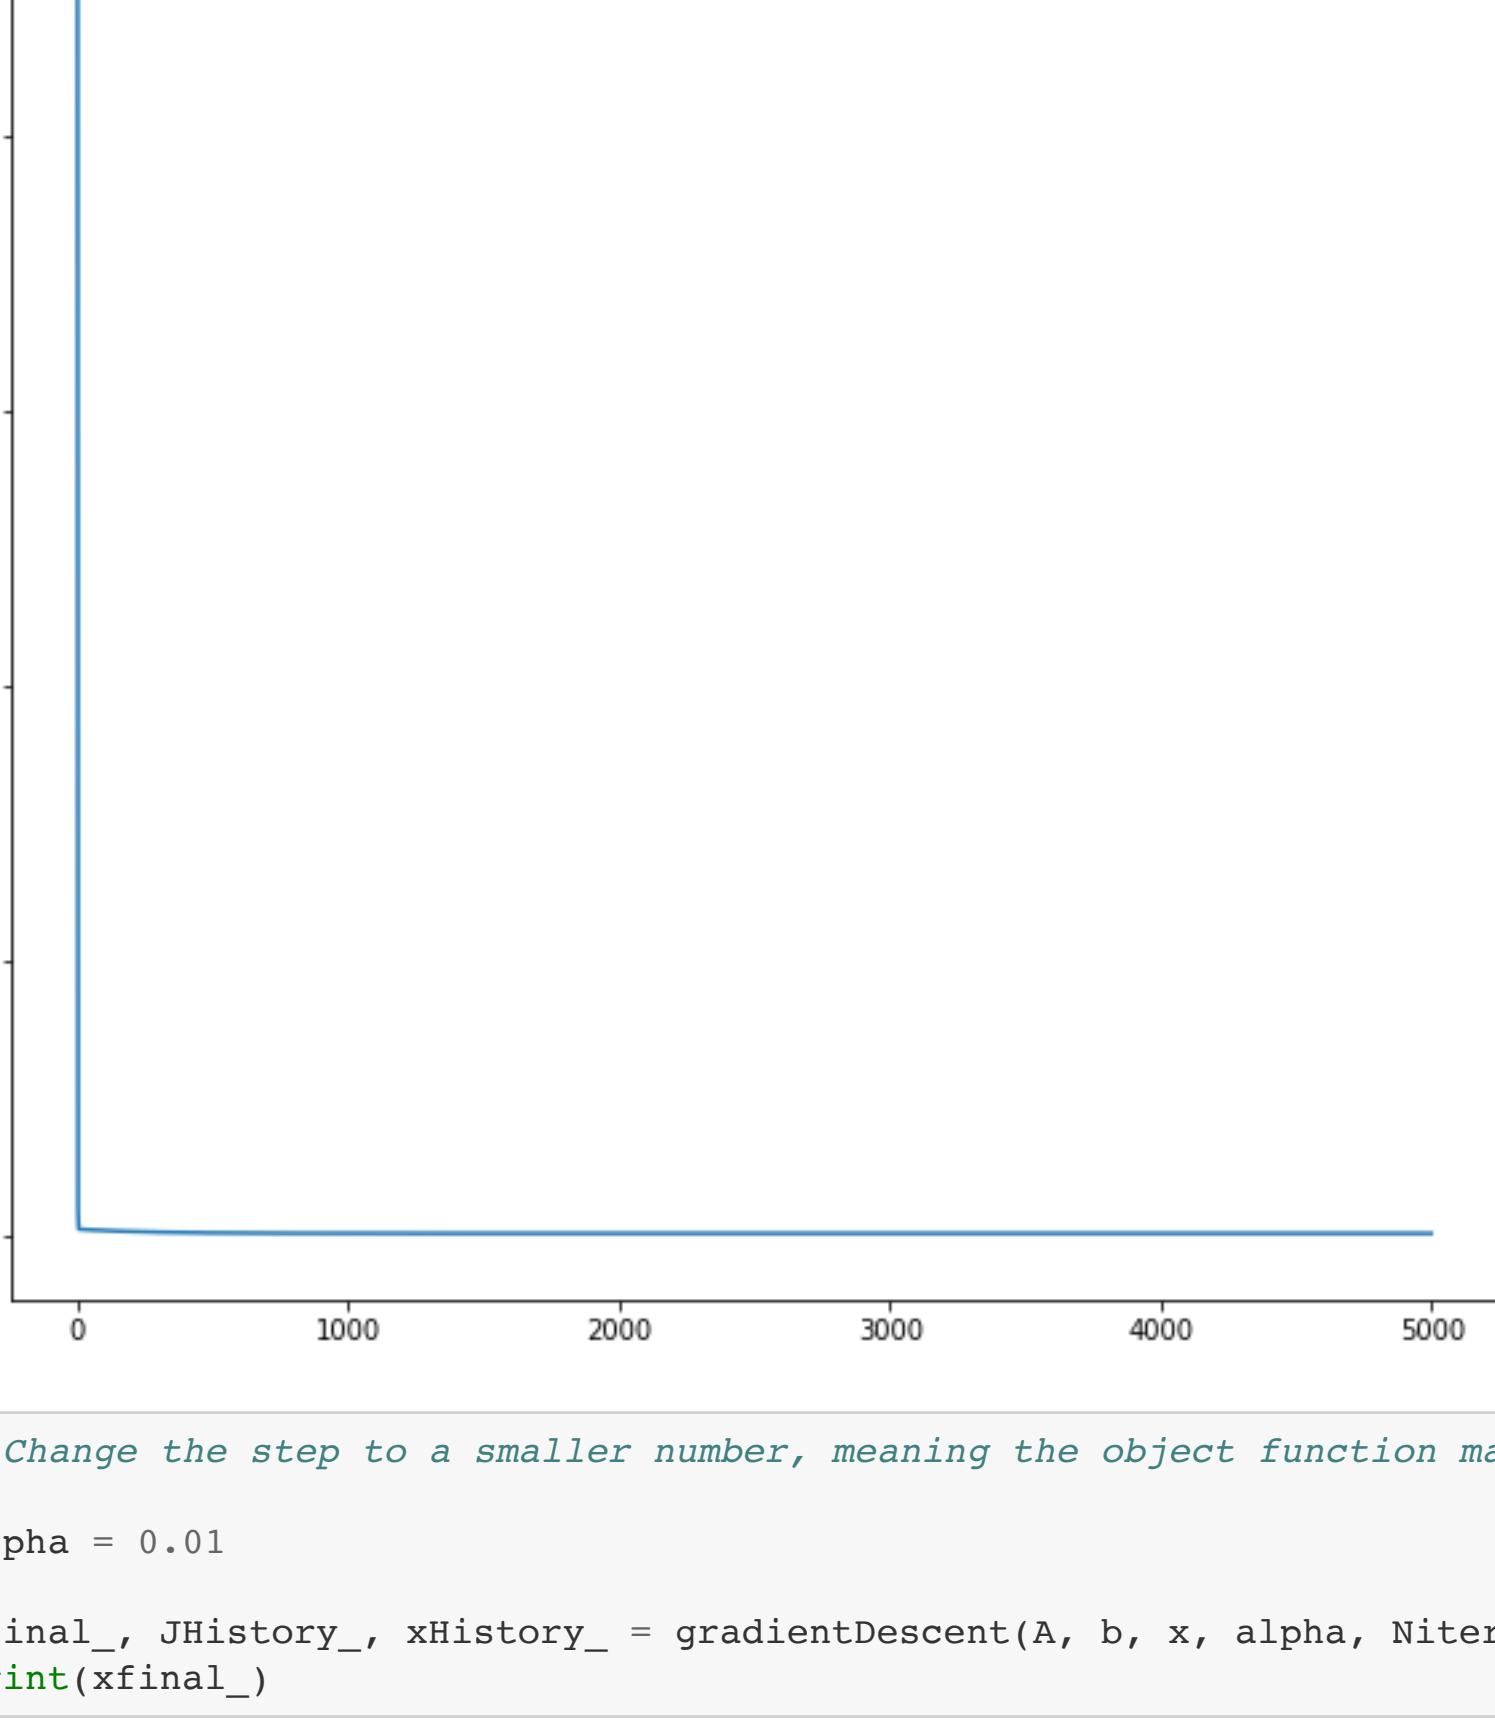
plt.figure(figsize=(10,10))
l1 = plt.plot(list(range(Niter)),list(xHistory[:,0]), label = "x[0]")
l2 = plt.plot(list(range(Niter)),list(xHistory[:,1]), label = "x[1]")
plt.legend()
plt.show()
```



```
In [ ]: # Plot the object function changes as the iterations increase.

plt.figure(figsize=(10,10))
plt.plot(list(range(Niter)),list(JHistory[:,0]))

Out[ ]: <matplotlib.lines.Line2D at 0x7f0203241240>
```



```
In [6]: # Change the step to a smaller number, meaning the object function may converge slowly.


alpha = 0.01
xfinal_, JHistory_, xHistory_ = gradientDescent(A, b, x, alpha, Niter)
print(xfinal_)

[[0.80067302]
 [1.77051845]]
```

The result is not as good as that when using a larger step.

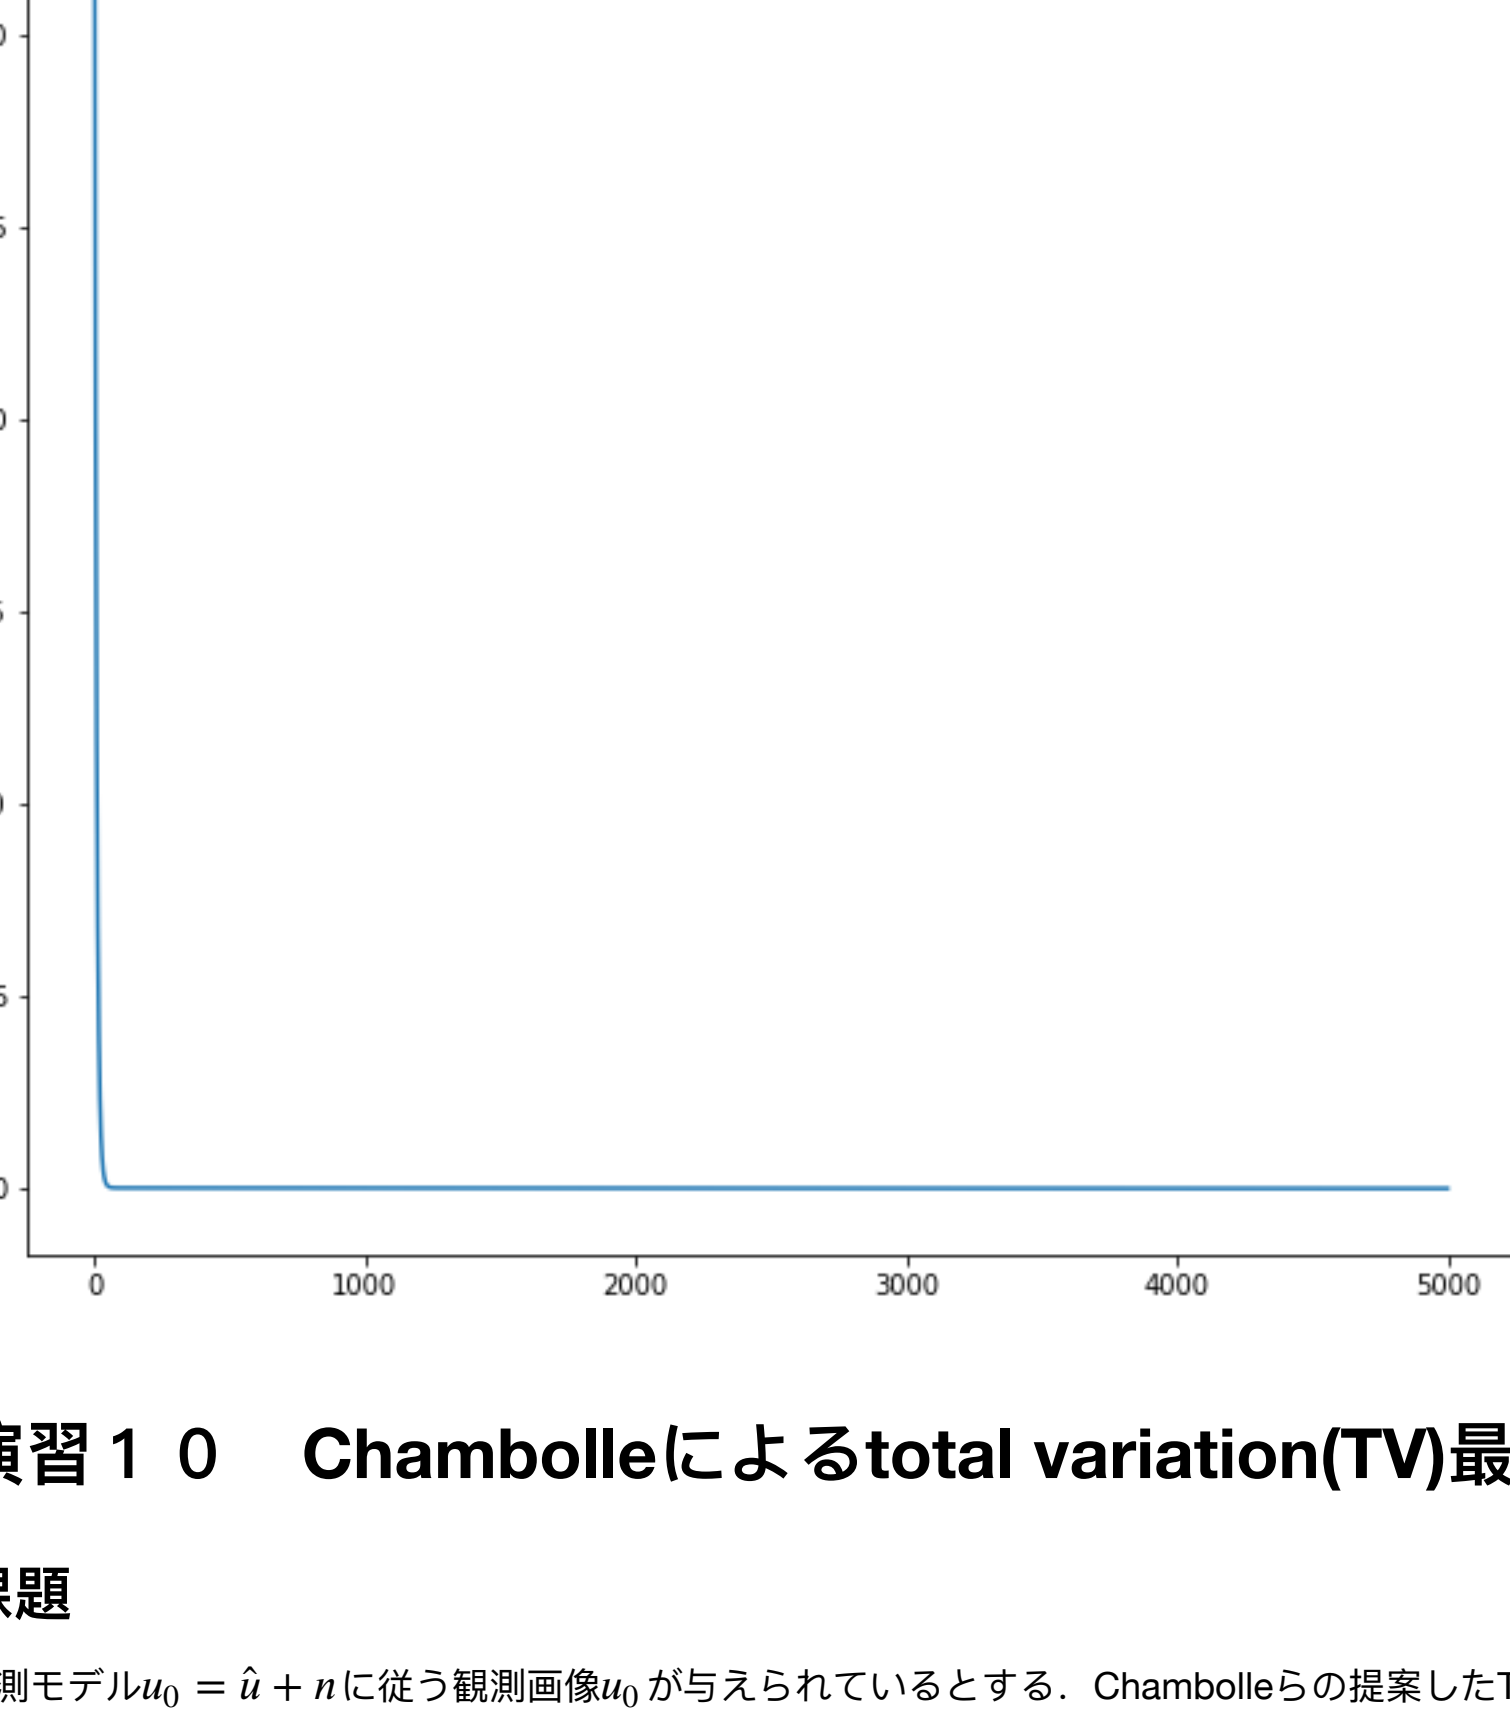
```
In [8]: #Plot of x[0], x[1]

plt.figure(figsize=(10,10))
l3 = plt.plot(list(range(Niter)),list(xHistory_[:,0]), label = "x[0]")
l4 = plt.plot(list(range(Niter)),list(xHistory_[:,1]), label = "x[1]")
plt.legend()
plt.show()
```



```
In [ ]: plt.figure(figsize=(10,10))
plt.plot(list(range(Niter)),list(JHistory_[:,0]))

Out[ ]: <matplotlib.lines.Line2D at 0x7f386d0ffda0>
```



演習 1 0 Chambolleによるtotal variation(TV)最小化アルゴリズムの実装

課題

観測モデル $u_0 = \hat{u} + n$ に従う観測画像 u_0 が与えられているとする。Chambolleらの提案したTV最小化の数値計算アルゴリズムにより、ノイズ除去を行い (u_0 に近いことが期待できる) 推定画像 u を求める手法を実装せよ。

Image denoising is a typical inverse problem, of which the selection of regularization terms is very important. For a unknown image \hat{u} , the observed model is u_0 :

$$u_0 = K\hat{u} + n$$

Total Variation is a regularization introduced by Rudin, Osher, and the following formula is called as ROF model.

$$u^* = \arg \min_u \frac{1}{2\lambda} \|u_0 - u\|_2^2 + \|Du\|_{1,2}$$

$D: \mathbb{R}^N \rightarrow \mathbb{R}^{2N}$ is the discrete gradient operator for position (i, j) in the image. And $\|Du\|_{1,2}$ is called as TV norm. In this view, the image denoising is performed as an infinite-dimensional minimization problem.

Chambolle proposed a dual-formulation-based projection algorithm for TV-denoising. The ROF model can be represented as

$$\min_u \frac{1}{2\lambda} \|u_0 - u\|_2^2 + \max_p \langle D^T p, u \rangle$$

where the $D^T: \mathbb{R}^{2N} \rightarrow \mathbb{R}^N$ is the conjugate operator that forms the following equation:

$$\langle p, Du \rangle = \langle D^T p, u \rangle$$

To simplify, with gradient direction $-D(D^T p - \frac{u_0}{\lambda})$, the pixels in the denoised image can be obtained using projected gradient method:

$$p_{i,j}^{k+1} = \frac{p_{i,j}^k - \tau((D(D^T p^k - \frac{u_0}{\lambda})))_{i,j}}{\max[1, |p_{i,j}^k - \tau((D(D^T p^k - \frac{u_0}{\lambda})))_{i,j}|]}$$

where $|z| := \sqrt{(z_1)^2 + (z_2)^2}$

The denoised image is obtained by:

$$\hat{u} = u_0 - \lambda D^T p$$

```
In [10]: !git clone https://github.com/mdipcit/standard_images/

Cloning into 'standard_images'...
remote: Enumerating objects: 10, done.
remote: Counting objects: 100% (10/10), done.
remote: Compressing objects: 100% (8/8), done.
remote: Total 10 (delta 0), reused 0 (delta 0), pack-reused 0
Unpacking objects: 100% (10/10), done.
```

```
In [11]: cd standard_images/

/content/standard_images
```

```
In [12]: import numpy as np
import matplotlib.pyplot as plt
from PIL import Image
from numpy import *
```

Define operator D and D^T .

```
In [13]: def D(im):
    h, w = im.shape
    dh = np.concatenate([im[1:, :] - im[0:-1, :],
                          np.zeros([1, w]),
                          np.concatenate([im[:, 1:] - im[:, 0:-1],
                                          np.zeros([h, 1]),
                                          axis=1])],
                          axis=1)
    return np.dstack([dh, dw])

def Dt(p):
    h, w_2 = p.shape
    return np.concatenate([-p[0, :, 0].reshape([1, w]), -p[1:-1, :, 0] \
                          + p[0:-2, :, 0], p[-2, :, 0].reshape([1, w]), axis=0] \
                          + np.concatenate([-p[:, 0, 1].reshape([h, 1]), -p[:, 1:-1, 1] \
                          + p[:, 0:-2, 1], p[:, -2, 1].reshape([h, 1]), axis=1])
```

Define the calculation of signal-to-noise ratio. High signal-to-noise ratio will be obtained if denoising is well done.

```
In [14]: def PSNR(im1, im2):
    return 10*np.log10(1/np.mean((im1-im2)**2))
```

```
In [15]: uh = np.asarray(Image.open('pepper.png')).astype(float)/255.0
h, w = uh.shape
sig = 10/255
u0 = uh + np.random.randn(h, w) * sig
print(PSNR(uh, u0))

28.100653200424873
```

```
In [16]: # initialization
tau=0.9/4 #反復法の勾配降下のパラメータτ
lamb = 0.03 #目的関数 (式(15))中の正則化項の重みλ
Niter = 100 #反復回数
# 反復中で使う変数を初期化
p = np.zeros([h, w, 2])

# Make a list of error to record the difference of updated U and previous U
error = list(range(Niter))

# The parameters include the observed image, the initialized p,
# the initial U, the gradient parameter tau, the weight of regularization,
# and the number of iterations
def denoise(im, p, U_init, tau, lamb, Niter):
    h, w = im.shape

    U = U_init
    Px = p[:, :, 0]
    Py = p[:, :, 1]

    for i in range(Niter):
        # Chambolle's iteration
        #for i in range(Niter):
        # ここにUの更新の式を書く
        Uold = U

        # The gradient of rows and columns is calculated separately
        GradUx = D(Dt(p) - im / lamb)[ :, :, 0]
        GradUy = D(Dt(p) - im / lamb)[ :, :, 1]

        PxNew = Px - tau * GradUx
        PyNew = Py - tau * GradUy

        NormNew = maximum(1, sqrt(PxNew**2+PyNew**2))

        Px = PxNew/NormNew
        Py = PyNew/NormNew

        p = np.dstack([Px, Py])

        #---dual-primal
        #u = u0 - lamb*Dt(p)
        U = im - lamb * Dt(p)

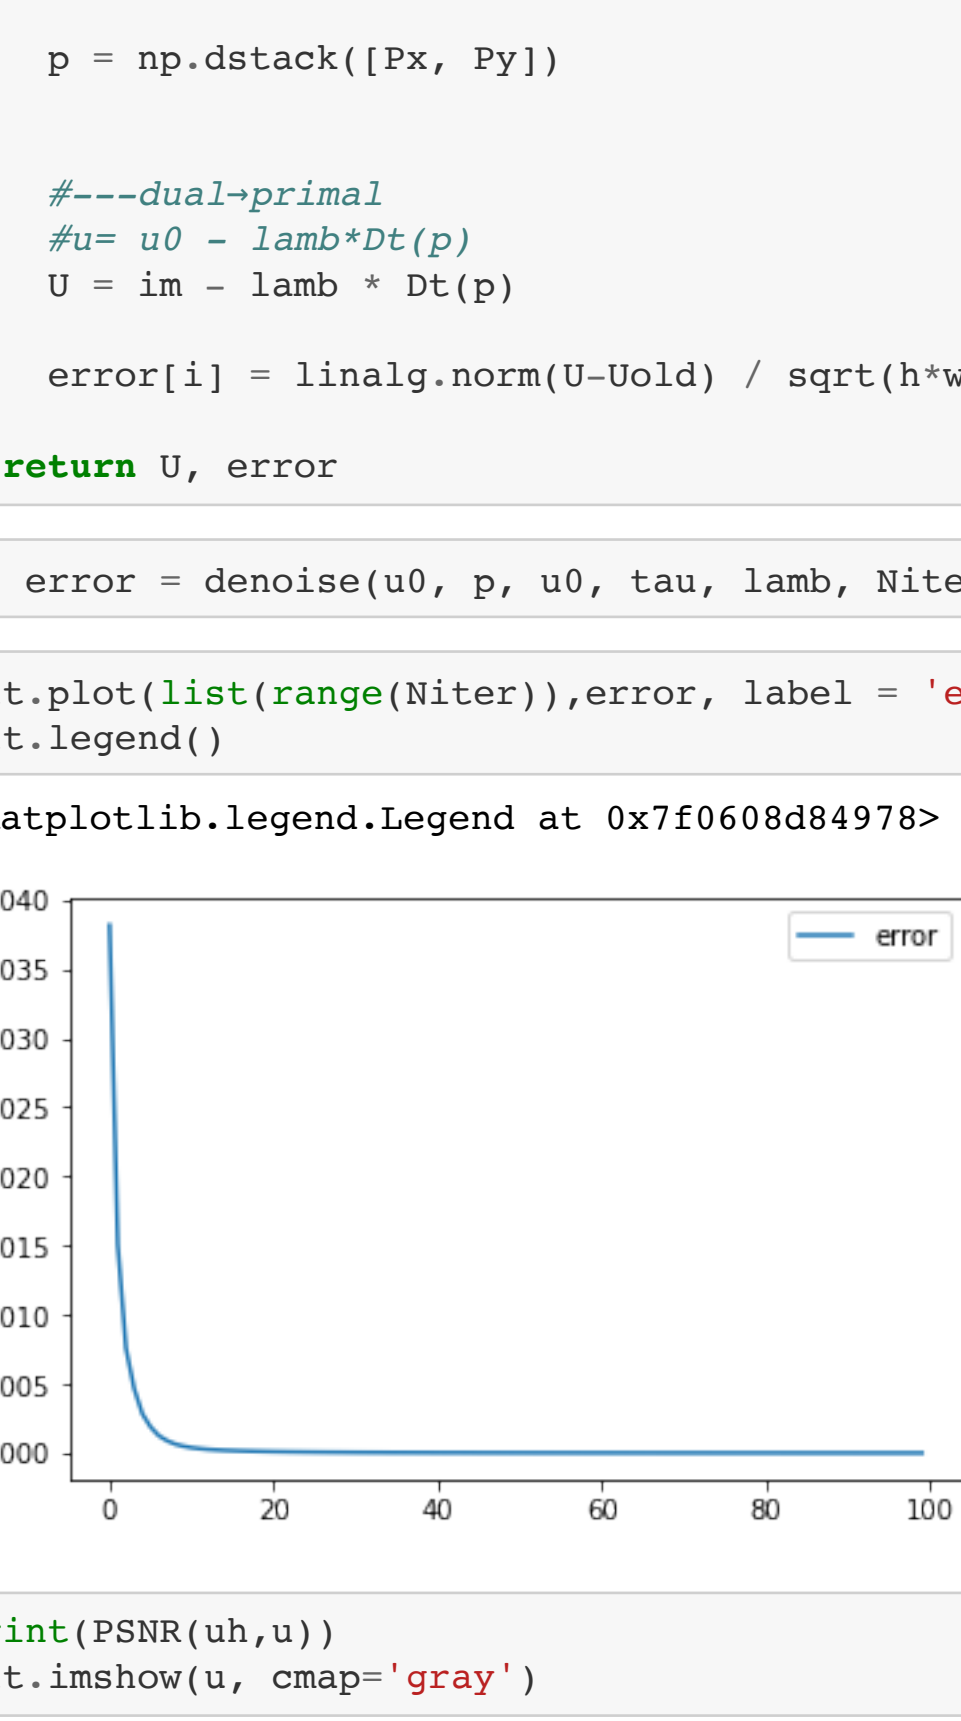
        error[i] = linalg.norm(U-Uold) / sqrt(h*w)

    return U, error
```

```
In [17]: u, error = denoise(u0, p, u0, tau, lamb, Niter)
```


```
In [18]: plt.plot(list(range(Niter)),error, label = 'error')
plt.legend()
```

```
Out[18]: <matplotlib.legend.Legend at 0x7f0608d84978>
```

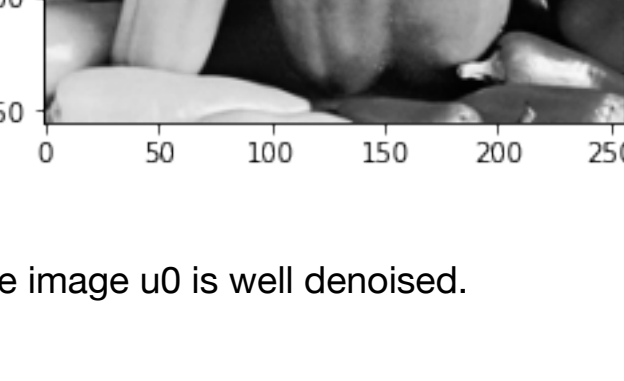


```
In [19]: print(PSNR(uh,u))
plt.imshow(u, cmap='gray')
```

```
Out[19]: 33.42813939954505
<matplotlib.image.AxesImage at 0x7fb3e0184a8>
```



```
In [20]: plt.imshow(uh, cmap='gray')
Out[20]: <matplotlib.image.AxesImage at 0x7f0605d5fda0>
```



The image u_0 is well denoised.